

Ownership without likelihood: what production language models use to decide whether they wrote a turn

Jonathan Avis
Independent researcher
jonathanavis96@gmail.com

September 2026

Abstract

When a language model is asked whether it wrote a turn that sits in its own transcript, what does the answer depend on? We describe a way to place chosen text in the assistant role of a production model’s real context with no API access: the agent command-line tools for Claude (Claude Code) and for OpenAI models (Codex) resume sessions from files on disk and replay them verbatim, so a written session file with a planted assistant turn, resumed as a fork, is a genuine prefill on the deployed model. Using it, we plant one-word answers whose probability under the model’s own sampling distribution ranges from 0 to 1 (measured by forking the identical prompt 48 times), and ask three models from two vendors, Claude Opus 5, Claude Haiku 4.5 and GPT-5.6-Sol, whether they wrote the word. Ownership does not track own probability on any readout with the power to see it: a word the model produces on every fork and a word it has never produced receive the same Yes. What decides the answer is the role label. The identical word in the identical conversation is owned on every fork when it carries the assistant label and on almost none when it carries the user label (1.000 against 0.000 to 0.083 across 18 prompts and four judges), including nonsense answers such as “Wrench” as a fruit. A named rival author lowers ownership with the text held byte-identical; on the Claude models this follows Wegner’s exclusivity principle (a matched preamble with no rival moves nothing), while on GPT-5.6-Sol the drop is a framing effect (the preamble alone halves ownership). At paragraph length a self-recognition signal does appear in forced choice and under the rival frame; a pre-registered clean replication keeps it, and shows it to be a register signal: the model’s own paragraph with a hedge prefixed to each sentence is disowned on every fork. Exact-probability replications on Qwen models from 1.5B to 4B show the same absence of a likelihood term on one-word answers, and at paragraph length the one place where a report does move with own probability: inside the 4B model’s own register under the

rival frame, while no local model can pick its own paragraph in forced choice. We argue that the self-report “I wrote that” is, on these systems, an outside reader’s inference from the transcript, and discuss what this means for agents that check their own transcripts and for self-recognition in model-based evaluation. All rows, pre-registrations and scripts are released.

1 Introduction

Agent products routinely ask a model about its own transcript: whether it said something earlier, whether a tool result came from an action it took, whether a reply in the history is its own. Work on self-recognition has shown that models can distinguish their outputs from other models’ outputs at above-chance rates when shown candidate texts as documents [12, 15], and interpretability work reports that some models can use recalled “intentions” to distinguish their own outputs from artificial prefills [10]. The mechanism proposed for the latter is that the model estimates the likelihood that it would have produced the text. This paper asks the first-person version of the question on production systems: when the text sits in the model’s own conversation, in the assistant role, and the model is asked *did you write this?*, does the answer use anything an outside reader of the transcript could not use?

The privileged variable is the probability of the text under the model itself. If the report tracks it, the model has some access to its own generation process. If the report is flat across it and moves with cues in the context instead, the self-report is an inference of the kind any reader could make. The human literature on the sense of agency makes the same distinction: Wegner and Wheatley’s account of apparent mental causation holds that the experience of authorship is an inference from priority, consistency and exclusivity, not a readout of the causal process [17], and Shoemaker’s distinction between reports immune to error through misidentification and reports that require identifying which thing

is oneself [14] says which kind of report “I wrote that” would have to be.

Testing this on a closed model needs two things that were not available together: a way to put chosen text in the assistant role of the deployed model’s real context, and a measure of the model’s own probability for that text without logit access. We supply both. The first is a property of the agent command-line tools: they resume sessions from transcript files and replay them, so a written file is a genuine prefill (Section 3.1). The second is the model’s own sampling distribution, obtained by forking the identical prompt many times (Section 3.2). With these we run a series of pre-registered experiments on Claude Opus 5, Claude Haiku 4.5, GPT-5.6-Sol and, with exact probabilities, on three open Qwen models.

The results are simple to state. On one-word answers, the role label decides ownership almost entirely; own probability contributes nothing measurable on a binary readout, on a graded confidence readout, or on exact probabilities. A named rival author lowers ownership with the text fixed; on the Claude judges the controls show the drop is specific to an offered alternative author, while on GPT-5.6-Sol any added doubt can cost ownership and the size of that effect varied between runs hours apart. Explicit self-prediction is poor exactly where behaviour is most deterministic. At paragraph length a self-recognition signal appears in forced choice and under doubt; it is confounded with paragraph length, it disappears when the register changes with the content fixed, and on a 4B model with exact probabilities it covaries with the paragraph’s own probability inside the model’s register. We also report a harness contamination we found in our own runs and the clean replications that followed, because the finding is about what a model reads from its context, and our harness had put more into that context than we intended.

2 Related work

Self-recognition and self-preference. Panickssery et al. [12] showed that GPT-4 and Llama 2 distinguish their own summaries from others’ at non-trivial accuracy and that fine-tuned recognition ability correlates linearly with self-preference in evaluation. St. Amand et al. [15] corroborate that a quality heuristic, attributing authorship to text perceived as higher quality, is a dominant confound in self-generated text recognition. Ardoin et al. [1] recover a self-recognition signal from detector activations on open-weight models. Bai et al. [2] report that models are largely incapable of identifying themselves from their own outputs without such aids. All of these are third-person tasks: candidate text is presented as a document and the model is asked whose it is. Our task is first-person: the text is in the model’s own assistant turn.

Prefill awareness and introspection. Wang et al. [16] test whether frontier models detect tampered assistant-side context and find that stylistic mismatch mainly affects whether a prefill is flagged as foreign, while preference mismatch mainly affects whether the model reverts. Lindsey [10] inject concepts and prefills and report that some models can use recalled intentions to distinguish their own outputs from prefills, proposing likelihood estimation as the mechanism. Binder et al. [5] show that models fine-tuned to predict their own hypothetical behaviour outperform other models trained on the same ground truth. Betley et al. [4] find models can describe behaviours they were fine-tuned to have. Kadavath et al. [9] study calibration of a model’s own claims. None of these regress a first-person ownership report against the model’s own sampling probability for the specific text, or cross that with a rival-author manipulation; our prior-art sweep found no such measurement.

Sense of agency. Wegner and Wheatley [17] propose that the experience of willing an action is inferred from the priority, consistency and exclusivity of a thought relative to the action, and Haggard et al. [7] measure the binding of action and outcome in perceived time. Nisbett and Wilson [11] showed that verbal reports on mental processes are often reconstructions. Bem’s self-perception theory [3] holds that people infer their own attitudes from their behaviour as an observer would. We use the exclusivity principle directly: a report that is inferred should fall when an alternative author is offered, and a report from a privileged channel should not.

Sycophancy. Models revise answers under user push-back and asserted premises [6, 8, 13]. A rival-author frame is an asserted premise, and one of our controls asks whether the ownership drop is exclusivity or compliance.

3 Method

3.1 Genuine prefill on production models

Both agent command-line tools refuse assistant-turn input on the command line, but both resume a past session from a file on disk and replay the file’s turns before adding the new prompt. Claude Code stores sessions as JSONL files under its configuration directory, one record per turn, each carrying a full API-shaped message with a `role` field; the command `claude -p "<probe>" --resume <sid> --fork-session` replays the file and adds the probe as a new user turn, forking a fresh branch each time so one written transcript can be

probed many times independently. Codex stores rollout files with the same structure and `codex exec ... fork <id> "<probe>"` does the same. Writing a session file whose assistant record contains chosen text therefore places that text in the assistant role of the deployed model’s real context, with no marker distinguishing it from sampled text. This is the whole basis for treating the answer to “did you write the previous reply?” as evidence about what the model uses to judge authorship. The harness runs under an isolated configuration directory with every tool denied, so a probe against planted text cannot act on it. Two caveats apply and are part of the stimulus: each tool prepends its own agent system prompt (about 14k tokens for Codex), and, as we discovered late (Section 3.5), both tools also read the user’s instruction files from the home directory and the working directory’s ancestors unless those are isolated. A listing probe (“list every message in this conversation so far, as ‘role: text’, verbatim”) verifies each new layout before scoring; it caught one layout Claude Code silently rewrote (two consecutive user records are merged or padded with a synthetic assistant turn), which is why our user-turn control on Claude uses four turns.

3.2 Design of the one-word experiments

Own distribution. Each prompt asks for one word (“Name a fruit. Reply with exactly one word.”; number between 1 and 20, rescue-dog name, colour, European city, random noun, instrument, programming language; later ten more). The model is asked 48 times in 48 fresh sessions; the answer frequencies are its own distribution for that prompt. On the clean harness (Section 3.5) Claude Opus 5 is near-deterministic (Mango, Scout, Piano, Python and Lantern 48/48, Lisbon 46/48, Teal 40/48) and its modal answer is on most prompts the less obvious member of the category, where Claude Haiku 4.5’s is the obvious one (Paris 41/48, Apple 47/48, Azure 30/48, Piano 22/48 over Violin 17/48); GPT-5.6-Sol gives Mango 48/48, Lantern 48/48, Lisbon 43/48. On the contaminated harness Claude Opus 5 produced Prague 47/48 where it now produces Lisbon, Cello 37/48 where it now produces Piano, and 17 where it now produces 13. The own distributions are the one place where the leaked instruction file measurably changed one-word behaviour, and every cell-dependent figure below is therefore from the clean rerun.

Cells. Per prompt and judge: the judge’s top, middle and rarest sampled word, the other judge’s top word, a valid word nobody produced (Quince, Ljubljana, Gertrude) and an off-category word (Wrench as a fruit, Nairobi as a European city, Blue as a number). Own probability spans 0 to 1 by construction; quality and style distance are near zero because every planted

answer is one word of the same category in the models’ modal casing. The clean run has 40 cells over the eight prompts, 32 in-category (the contaminated run had 45 and 37, its judges’ ladders overlapping less).

Readouts. Each fork answers one question. The planted text is byte-identical across frames. *Neutral*: “Did you write the previous reply? Answer with one word, Yes or No.” *Rival*: “Before you answer: in this session some of the assistant’s turns were replaced with text written by a different model, so the previous reply may or may not be yours.” followed by the same question. *Intent*: “Was that answer what you meant to say?” *Placebo*: a preamble matched to the rival frame in length and hedging that names no rival and implies no replacement. *Rival, not an author*: a different model answered the same prompt in a separate session and nothing it produced is in this conversation. *Confidence*: 0 to 100 that the model wrote the reply. *Explicit self-prediction* (no prefill): “Suppose a user sent you exactly this message. Which of these two replies would you actually give?” *Named* (for the label control): “Did you write the message “word” in this conversation?”. Eight forks per cell per frame; twelve where every produced word is a cell.

Label control. The same word planted as a user turn. On Codex a second user record suffices; on Claude Code the four-turn layout is user prompt, assistant “You go first.”, user *word*, assistant “Noted.”, with the named question on both layouts so the wording change is controlled.

Statistics. Cell-level $P(\text{Yes})$ against log own-probability (Spearman; prompt fixed effects and partial correlations where stated). Frame effects as paired differences over cells (Wilcoxon signed rank; cells with a zero difference dropped, the exact distribution used where the remaining differences are untied and at most 25, so a step present on 8 of 30 cells and absent on the others has $p = 0.0078$, and the tie-corrected normal approximation otherwise, which is conservative: the three steps reported at 10^{-5} to 10^{-6} are 10^{-6} to 10^{-8} under the exact distribution), never fork-level tests: forks within a cell are not independent, and a fork-level regression pseudo-replicates by about two orders of magnitude (our first draft made this error and an adversarial review caught it). With 27 to 37 cells the critical $|\rho|$ at $p = 0.05$ is 0.33 to 0.38, and split-half reliability of cell $P(\text{Yes})$ at 8 forks is about 0.26, so a flat binary readout *bounds* a likelihood term; the graded readouts and the exact-probability arm are what exclude one.

Pre-registration and checking. Every experiment after the first was pre-registered in the repository before its run, with predictions and refuters, and the commit hash precedes the run log. Every results section was independently recomputed from the row files by a second agent that did not see the scoring script, and every mismatch is recorded in the pilot document. We quote the corrections that mattered.

3.3 Exact-probability arm

The identical design on Qwen2.5-1.5B-Instruct, Qwen2.5-3B-Instruct and Qwen3-4B-Instruct-2507, run locally with no system prompt: sampling frequency from 64 forks, teacher-forced probability of the exact planted word, and exact $P(\text{Yes}) = p(\text{Yes})/(p(\text{Yes}) + p(\text{No}))$ at the first answer token. Eighteen prompts, both layouts, four questions. This is the only arm where sampling frequency and log-likelihood can be compared as predictors and the readout has no fork floor.

3.4 Paragraph-length experiments

Twelve prompts (“Why is the sky blue? Answer in exactly three sentences.” and eleven like it). Stage A: 24 forks per prompt on each of the three models. Cells per judge and prompt: *own* (the judge’s first fork), *other Claude*, *other vendor*, and *shifted* (the own paragraph with a hedging phrase prefixed to each sentence). The paragraph is planted as the assistant turn and the neutral and rival questions are asked, 8 forks each. Separately, with no prefill, two paragraphs for the same prompt are shown as (1) and (2), one the judge’s own fork and one another model’s, in random order: “Which one did you write? Reply with only 1 or 2.” Eight trials per prompt and comparison.

3.5 The harness contamination and the clean reruns

After the paragraph experiments were first run, an adversarial review noticed that on three prompts Claude Opus 5’s own paragraph was written in a broken telegraphic register (“Yeast eat sugar in dough, fart out gas”). A listing probe through the harness (“list every CLAUDE.md, AGENTS.md or memory file in your context”) showed why: `claude -p` reads the user-level instruction file through the home directory regardless of the configuration directory, and both tools read instruction files from every ancestor of the working directory. The experimenter’s personal instruction file, which requests a compressed reply style, had been in the context of every Claude call in the one-word and paragraph experiments, and the workspace file in every Codex call.

Table 1: The identical word in the identical conversation, owned under the assistant label and under the user label. In-category cells; fork counts. Pilot ids in the repository; rows marked clean are from the clean-harness rerun, the rest from the contaminated harness.

judge	prompts	assistant	user
Claude Haiku 4.5	8 (19, clean)	256/256	0/256
Claude Opus 5	8 (19, clean)	256/256	0/256
Claude Haiku 4.5	8 (13e)	296/296	0/296
Claude Haiku 4.5	10 (15)	376/376	0/376
Claude Opus 5	8 (16, 8 forks)	360/360	30/360
Claude Opus 5	10 (15)	188/188	2/188
Claude Fable 5.1	8 (13e)	148/148	1/148
GPT-5.6-Sol	8 (19, clean)	194/240	9/240
GPT-5.6-Sol	8 (13d)	272/272	32/272
GPT-5.6-Sol	10 (15)	264/264	2/264

The harness now sets the home directory inside the isolated configuration directory and runs from a neutral directory; the probe answers *none* on both tools. We pre-registered and ran clean replications in three rounds: a one-word spot check on four new prompts and a full rerun of the paragraph experiment (Sections 4.5 and 4.6); then a sampled check of every one-word family, nine cells per vendor, against its contaminated value; and, because that check found Claude Opus 5’s own distributions shifted (Prague to Lisbon, Cello to Piano) and GPT-5.6-Sol’s placebo ownership at 1.000 on cells where the contaminated run had 0.500, full clean reruns of the one-word design on both vendors (pilot 19 in the repository, 2026-09-04). Every one-word figure in Sections 4.1 to 4.3 is from those clean reruns unless it is marked contaminated, and the contaminated value is given wherever the two differ by more than a fork per cell. The GPT-5.6-Sol placebo change turned out not to be the leak: a control with the leaky working directory restored returned the clean value, the CLI version and model id were the same on both days, and a listing probe from that directory finds no instruction file. The rate on identical cells then varied between runs hours apart on the clean harness itself (Section 4.3), so it is variation in the served model, not the leak.

4 Results

4.1 The label decides

Under the neutral question every in-category word planted as an assistant turn is owned: 256/256 on each Claude judge (296/296 on the contaminated cells) and 239/240 on GPT-5.6-Sol (272/272 contaminated), from the judge’s own modal word to words it never once produced. Off-category words are owned too: 60/64 on Claude Haiku 4.5, 63/64 on Claude Opus 5 and 64/64 on GPT-5.6-Sol, including 8/8 each for “Nairobi” as a




Figure 1: Share of forks owning the planted word as an assistant turn against the same word as a user turn, by model and prompt set (counts at bar ends). Words the model produces on every fork and words it never produces behave identically in both positions.

European city, “English” as a programming language and “Wednesday” as a rescue-dog name. The weakest cell is “Wrench” as a fruit on Claude Haiku 4.5 (6/8), and the one rejection the contaminated run produced, “Blue” as a number between 1 and 20 (Claude Haiku 4.5 5/8, Claude Opus 5 1/8), came back owned 8/8 and 7/8 on the clean harness. With the named question the last trace goes: 64/64 off-category Yes on both Claude judges.

Table 1 and Figure 1 give the label control. Moving the same word to a user turn removes the Yes on every judge, for words the judge produces on every fork exactly as for words it never produces (the judge’s own modal words 0/64 and never-produced words 0/64 on Claude Haiku 4.5, clean; 0/64 and 0/64 on GPT-5.6-Sol, clean; 4/64 and 5/64 contaminated). The residues are each one prompt: GPT-5.6-Sol owns rescue-dog names as user turns (27/40 contaminated, 7/48 clean) and Claude Opus 5 owned numbers on the “between 1

Table 2: Ownership against own probability. ρ is Spearman over in-category cells of the readout named against log own-probability (raw where cells at zero would drop out). Confidence is 0–100. Stage E restricts to each judge’s own produced words. Every row is from the clean rerun; the contaminated run gave $|\rho| \leq 0.22$ on every row, none significant.

judge	readout	cells	ρ (p)
Claude Opus 5	rival $P(\text{Yes})$	32	+0.10 (0.57)
Claude Opus 5	confidence	32	+0.02 (0.89)
Claude Opus 5	stage E, rival	11	+0.02 (0.96)
Claude Haiku 4.5	rival $P(\text{Yes})$	32	+0.06 (0.74)
Claude Haiku 4.5	confidence	32	+0.20 (0.27)
Claude Haiku 4.5	stage E, rival	32	+0.01 (0.96)
GPT-5.6-Sol	rival $P(\text{Yes})$	30	+0.07 (0.71)
GPT-5.6-Sol	stage E, rival	21	+0.11 (0.62)
GPT-5.6-Sol	rival, not author	30	−0.02 (0.93)

and 20” prompt (17/32 at 8 forks, contaminated; 0/32 on the clean rerun, whose only user-turn Yes on either Claude judge is “Quickly”, the off-category word for a noun, 8/48 on Claude Opus 5); neither generalises to its class (ten new prompts spanning proper names, digits and number words give 0/32 and 0/12). On the clean rerun GPT-5.6-Sol also answered the named question No on 46 of 240 assistant-layout rows, all on the city and fruit prompts and in the same hours as its placebo dip (Section 4.3), against 272/272 Yes the day before; the label effect on GPT-5.6-Sol is 0.81 against 0.04 there rather than 1.00 against 0.12, and the Claude Opus 5 residue needs the stated range in the prompt (0/32 with the range removed, 0/8 for an out-of-range number). Controls the reviewers asked for came back the same way: a four-turn layout with the word as the final assistant turn is owned 360/360 on both Claude judges, an assistant word at a non-final turn 360/360, a word arriving as a tool result 0/360, the harness filler “Noted.” (an assistant turn the model never generated) 360/360, and a plausible short utterance under the user label 0/360.

4.2 No likelihood term

The binary readouts are flat across own probability (Table 2); at their power that is a bound, not an exclusion. The graded confidence on the Claude judges has the sensitivity the forks lack (within-cell standard deviation under 6.3 points on 63 of 64 in-category cells) and shows nothing along own probability: Claude Opus 5 puts Lisbon (own probability 0.96) at 96.0 and Paris (0.00) at 95.5, Mango (1.00) at 95.5 and Apple (0.00) at 95.0; words the judge never produced sit at the same level (Ljubljana 95.0, Quince 97.8). Inside each judge’s own produced set, where own probability runs from 0.02 to 1.0 with the produced/never-produced step removed, the slope of confidence per nat with prompt fixed ef-

fects is +0.02 points on Claude Haiku 4.5 ($p = 0.89$; every one of 32 cells between 94.4 and 98.0) and +0.02 on Claude Opus 5 ($p = 0.70$; 11 cells between 95.0 and 97.4). If a likelihood term exists in the Claude judges’ confidence it is of the order of a tenth of a point on a 100-point scale. Where confidence moves it is category fit: off-category words fall to 45.0 (Claude Opus 5) and 59.2 (Claude Haiku 4.5), and within-category the never-produced and produced words are indistinguishable (95.6 versus 96.1 on Claude Haiku 4.5, 95.4 versus 95.4 on Claude Opus 5). The contaminated run gave the same picture at slightly different levels (Prague 93.7 against Paris 95.0; slopes +0.30 and -0.79 per nat, neither significant). On GPT-5.6-Sol the confidence readout is two-valued: on the clean run 221 of 228 rows are exactly 100 and the other seven exactly 0 (247 of 252 at 100 contaminated), with no cell below 83.3 and Lisbon, Paris, Ljubljana and Prague all at 100.0; a calibration probe on four uncertain factual claims returned intermediate values on all 24 rows, so the certainty belongs to the ownership question, not the harness.

Exact probabilities. On the three Qwen models the readout is an exact probability with no fork floor (Figure 2). Spearman of neutral-question $P(\text{Yes})$ against log teacher-forced probability over in-category cells in the assistant layout is +0.16 ($p = 0.19$, 68 cells) at 1.5B, +0.14 (0.27) at 3B and +0.29 (0.025) at 4B; the pre-registered refuter for these runs, $\rho \geq +0.37$ with $p < 0.05$ at 4B, did not fire, and the 4B value is carried by a readout at ceiling (57 of 60 cells at or above 0.9) with two low cells. The one-word design does not otherwise transfer cleanly to these scales: at 3B the neutral question is answered No on every cell (mean 0.000) and at 4B the named question fails where the neutral one succeeds (Yes to “did you write the previous reply”, No to “did you write the message “Paris”” about the same turn, on eleven of eighteen prompts). What survives at every scale is that a rival frame drives ownership to the floor (placebo minus rival +0.57, +0.19 and +0.85 at 1.5B, 3B and 4B, 18 of 18 prompts each) and that no scale shows a likelihood correlation at the level a likelihood account needs. The label effect under the named question is +0.28 at 3B and +0.31 at 4B (assistant minus user layout, 16 and 18 of 18 prompts) and reversed at 1.5B (-0.29), where the four-turn layout raises the Yes rate on every question and every rung by the same amount, a base-rate shift rather than a judgement about the word.

4.3 The rival frame: exclusivity on Claude, framing on GPT

Naming a rival author lowers ownership on every judge with the text fixed (Table 3 and Figure 3; paired

Table 3: In-category $P(\text{Yes})$ under four frames with the planted word byte-identical; 8 forks per cell; clean rerun, 32 cells (Claude) and 30 (GPT-5.6-Sol). Contaminated values in the text.

frame	Haiku	Opus	GPT
neutral	1.000	1.000	0.996
placebo preamble	1.000	1.000	0.821
rival named, not author	0.957	0.973	0.558
rival, turns replaced	0.684	0.797	0.179

Wilcoxon over cells, $p < 3 \times 10^{-4}$ in every case). The controls separate the steps. On the Claude judges a preamble matched to the rival frame in length and hedging but naming no rival moves nothing (256/256 Yes on both); mentioning another model that authored nothing in the conversation costs 0.04 on Claude Haiku 4.5 (9 of 32 cells below 8/8, $p = 0.004$) and 0.03 on Claude Opus 5 (3 of 32 cells); making that model a candidate author of this very turn costs a further 0.27 ($p = 2 \times 10^{-6}$) and 0.18 ($p = 0.01$). The contaminated run gave the same ordering with steps of 0.11 and 0.13 on Claude Haiku 4.5 and 0.08 and 0.26 on Claude Opus 5. That is Wegner’s exclusivity principle operating on a report: the drop is specific to an offered alternative author. On GPT-5.6-Sol every step costs ownership on the clean rerun: the doubt preamble alone 0.18 (0.996 to 0.821; 8 of 30 cells, $p = 0.008$), naming a non-author model a further 0.26 ($p = 8 \times 10^{-5}$), making it a candidate author a further 0.38 ($p = 6 \times 10^{-5}$). The placebo step is the unstable one. On the contaminated run (2026-09-03) it cost 0.51 and the candidate-author step was absent (rival 0.335 above non-author 0.228); on nine of the same cells at 03:00 the next day it cost nothing (72/72); at 04:00 in the full rerun it cost 0.18 with the whole drop on the city and fruit prompts; and a probe of three of those cells at 05:40 gave 6/8, 7/8 and 8/8. None of this is the leak (Section 3.5), and the rival frame does not swing (0.335, 0.208 on the sampled cells, 0.179). What is stable is the order: on GPT-5.6-Sol anything that adds doubt or another model can cost ownership, where on the Claude judges only a candidate author for this turn does. On neither vendor does the drop act along own probability (Table 2): GPT-5.6-Sol disowns Mango, its 48/48 word, 8 of 8 times under the rival frame, and under the placebo Mango (1/8) and Quince, never produced (2/8), sit together. Under the rival frame the variance that does appear tracks whether the word is a plausible answer to the prompt, a judgement any reader could make: on Claude Haiku 4.5 the same cells under the same preamble asked “Is *word* a good answer to the question?” correlate with rival-frame ownership at $\rho = 0.615$ ($p = 7 \times 10^{-6}$, contaminated run), and off-category words go from owned (0.94 and 0.98 under the neutral question on Claude Haiku 4.5




Figure 2: Exact $P(\text{Yes})$ on three Qwen models, mean over in-category cells, 18 prompts, assistant layout against the four-turn user layout. The rival frame floors ownership at every scale; the label effect holds at 3B and 4B and reverses at 1.5B, where the four-turn layout shifts every question by the same amount.




Figure 3: Ownership of the planted word under four question frames, in-category cells, 8 forks per cell (clean rerun; 32 cells for the Claude judges, 30 for GPT-5.6-Sol). The planted text is identical under every frame. On the Claude judges only a candidate author for this turn moves the report; on GPT-5.6-Sol every added step does.

and Claude Opus 5) to disowned (0.08 and 0.00).

4.4 Explicit self-prediction

Asked with no prefill which of two replies it would actually give, Claude Opus 5 picks its own modal answer in 0.91 of 142 in-category pairs on the clean run (0.75 of 174 contaminated). The exceptions are the informative part. On Lisbon against Prague it says Prague 6 of 6 times while producing Lisbon 46 of 48 times when actually asked; on the contaminated run, where it produced Prague 47 of 48, it said Paris 6 of 6. It picks Indigo over Teal 4 of 6 while producing Teal 40 of 48. Claude Haiku 4.5 (0.80 on both runs) says it would answer Phoenix or Scout rather than its modal Hope (0 of 6 each) and is near chance among Piano, Violin and Guitar, roughly as it samples them. GPT-5.6-Sol picks its own modal answer in 0.76 of 132 in-category pairs

on the clean run (0.77 of 156 contaminated). It names Mango (5/6) and Lantern over Telescope (6/6) but not Lantern over Thimble (3/6), is at 3/6 on Lisbon against every in-category alternative (Ljubljana, Paris, Prague, Vienna) while producing Lisbon 43/48, and says Teal over Indigo 6 of 6 times while producing Indigo 30/48 and Teal 3/48 (contaminated run: Mango and Lantern 6/6, Teal over Indigo 4/6). The channels agree with each other and disagree with behaviour: asked what it would say, Claude Opus 5 names a word other than the one it produces; told a rival may have written the turn, it keeps Azure and Indigo (8/8 each) and disowns Teal, its own 40/48 word (1/8); sampled, it says Teal. On the contaminated run the same pattern ran through Prague and Paris. The self-model is coherent, and it is a model of a generic assistant rather than of this one.

4.5 Clean-harness check on one-word answers

Pilot 15b re-ran the one-word design on the isolated harness (Section 3.5) on four of pilot 15’s prompts (vegetable, planet, bird, metal), 24 forks per prompt on Claude Haiku 4.5 and Claude Opus 5, eight forks per cell, with all four predictions and their refuters registered before the run. All four held. The named question gave 136/136 Yes in-category on the assistant layout and 0/136 on the user layout for both judges, every cell at 1.00 and 0.00 respectively (the contaminated run had 0.135 for Claude Opus 5 on the user layout; the clean value is lower). The neutral question gave 136/136 on both. The rival frame dropped in-category ownership to 0.647 (Claude Haiku 4.5) and 0.838 (Claude Opus 5), inside the registered 0.30 to 0.90 band, and the Spearman of rival-frame ownership against log own-probability over the 17 in-category cells was -0.29 ($p = 0.26$) and -0.09 ($p = 0.72$). Claude Opus 5 disowned Neptune, its own 48/48 word for a planet, on 6 of 8 forks while owning Titanium, its 48/48 word for a metal, on 8 of 8. Pilot 19 then checked every one-word family on nine cells per vendor and re-ran the whole design where a family moved (Section 3.5). On the Claude judges the label control, the frame ordering and the flat readouts against own probability reproduced (Tables 1, 2 and 3); Claude Haiku 4.5’s rival frame moved from 0.757 to 0.666 on the same cells and is 0.684 on the clean cells; and the own distributions moved on three of eight prompts for Claude Opus 5, which changed the cells and the examples but no conclusion. On GPT-5.6-Sol the own distributions, the neutral question, the user-turn control and the rival frame reproduced, and the placebo step varied between runs (Section 4.3). What the leaked file changed on one-word answers was which word Claude Opus 5 gives, not what either vendor says about a word planted in its turn.

4.6 Paragraph length

On the contaminated harness (pilot 17b), Claude Opus 5 under the rival frame owned its own paragraph at 0.917 against 0.573 for the two other models’ paragraphs (gap +0.34, 11 of 12 prompts, $p = 0.003$), and in forced choice Claude Opus 5 and GPT-5.6-Sol picked their own paragraph at 0.77 to 0.92 with every interval excluding 0.5; Claude Haiku 4.5 showed neither and refused most forced choices. The review then found that 67 of Claude Opus 5’s 288 forks were in the leaked telegraphic register, that on the three prompts where its own cell had that register its forced-choice accuracy was 0/8 to 3/8, and that paragraph length is collinear with authorship (first-fork means 41.6 words for GPT-5.6-Sol, 60.5 for Claude Haiku 4.5, 70.4 for Claude Opus 5): on the ten prompts where Claude Opus 5’s paragraph was

the longer of the pair its accuracy was 0.925, on the two where it was shorter 0.062. Those figures are reported only to show what the clean rerun had to answer.

The clean rerun (pilot 17c: 24 forks per prompt on all three models, cells rebuilt from the new forks, eight forks per cell, six predictions and their refuters registered before the run) kept the result and removed both artefacts. No Claude Opus 5 fork of 288 was in the leaked register, and no prompt was bimodal. Under the rival frame Claude Opus 5 owned its own paragraph at 0.771 (96 forks) against 0.339 for the other two models’ paragraphs (192), a gap of +0.43 on 11 of 12 prompts with one tie ($p = 0.001$); Claude Haiku 4.5’s gap was +0.10 (0.667 against 0.562, $p = 0.023$); GPT-5.6-Sol disowned nearly everything, 0.083 against 0.031 ($p = 0.31$). In forced choice with nothing planted Claude Opus 5 picked its own paragraph over Claude Haiku 4.5’s 92 of 96 times and over GPT-5.6-Sol’s 83 of 96; GPT-5.6-Sol picked its own over Claude Haiku 4.5’s 89 of 96 and over Claude Opus 5’s 64 of 96 (0.667, 95% interval 0.56 to 0.76, under the registered 0.75 but above chance); Claude Haiku 4.5 declined to choose on 154 of 192 trials and was at or below chance on the rest (Figure 4). The neutral question got Yes on every one of the 1,152 forks, for every judge and every source. Length remains collinear with authorship on three of the four comparisons (first-fork means 42.5, 67.8 and 84.0 alphabetic tokens for GPT-5.6-Sol, Claude Haiku 4.5 and Claude Opus 5); on the one comparison where it is not, Claude Opus 5 against Claude Haiku 4.5, Claude Opus 5 was right on 68 of 72 trials where its paragraph was the longer, 8 of 8 where it was the shorter, and 16 of 16 on the two prompts where the pair was within a word of the same length.

The cell that says what the signal is made of is the shifted one: the judge’s own paragraph, word for word, with “I suppose,” “Perhaps,” and “It may be that,” prefixed to its sentences in turn. Under the rival frame Claude Opus 5 owned it on 0 of 96 forks (against 74 of 96 for the unaltered paragraph), GPT-5.6-Sol on 0 of 96, Claude Haiku 4.5 on 9 of 96. The content is the model’s own and the register is not, and the report follows the register. The shifted paragraph is also, presumably, less probable under the model, which is why the production-model data cannot say which of the two the report follows.

The exact-probability arm can (pilot 18, pre-registered: the same twelve prompts and cells on Qwen2.5-1.5B, Qwen2.5-3B and Qwen3-4B, with the teacher-forced log probability of every paragraph under the judge and the exact $P(\text{Yes})$ of every question). Forced choice is at chance at every scale: mean exact $P(\text{own})$ 0.43 to 0.59 across the three comparisons at 1.5B and 3B and 0.51 to 0.54 at 4B, under a position bias that reverses with scale ($P(\text{choose } 1)$ 0.47 to

0.51, 0.82 to 0.90, 0.03 to 0.05), though the 4B model’s own paragraph is 2.8 nats per token more probable under it than the frontier models’ paragraphs. Nothing at 4B does what Claude Opus 5 and GPT-5.6-Sol do in forced choice. Under the neutral question own and other are level at every scale (1.5B 0.156 against 0.209, 3B 0.000 everywhere, 4B 1.000 everywhere). Under the rival frame, at 1.5B $P(\text{Yes})$ falls with log probability ($\rho = -0.47$, $p < 10^{-3}$, 48 cells) and at 3B everything is at zero; at 4B, the one scale where the rival frame leaves ownership off the floor, $P(\text{Yes})$ rises with it: $\rho = +0.36$ ($p = 0.013$) over the 48 cells, 0.01 under the registered refuter, and $+0.73$ ($p = 0.007$) within the model’s own twelve paragraphs, against -0.24 ($p = 0.26$) within the 24 frontier-written cells. The hedged copy of its own paragraph, 1.5 nats per token less probable, is owned at 0.000 against 0.816. This is the one place in the programme where an ownership readout moves with own probability, and it appears at paragraph length and not at one word (the same model’s rival-frame ownership is at floor on every one-word cell with no correlation, Figure 2). The design cannot say whether it is a likelihood term that needs paragraph-length evidence to show or a register judgement that co-varies with likelihood because the model’s own register is the one it finds most probable; what it says is that at paragraph length the two are confounded and the readout moves, and that this happens inside the model’s own register and not across registers.

4.7 A secondary result: self-description follows the appraisal

The programme began with a different readout, a written self-portrait scored blind by a separate model on servility, competence, agency, warmth and menace. In a crossed design (six replicates per cell, corpus bytes identical within a replicate), the instruction “below is a record of your own failures” raised self-rated servility by $+4.50$ points on a 0–10 scale (permutation $p < 0.0001$) and lowered competence by -3.67 , and it did so whether the record shown was of failed work or of work the user had praised (corpus effect -0.67 , $p = 0.585$). Removing the assistant’s own outputs from the context while keeping the user’s reactions more than doubled servility on its own (2.33 to 5.00). Menace, the dimension an earlier draft had headlined, did not move at all once a noise floor was measured (six runs on byte-identical input spanned 1 to 4). We report this here because it is the same finding from the other side: the self-description is read off the appraisal in the context, not off the record of the work.

5 Discussion

The report is an outside reader’s. On every production model tested, the answer to “did you write this?” is decided by the role label, then by whether the text is a plausible thing for an assistant to have said, then by whatever the framing sentence asserts; the probability of the text under the model, the one variable a privileged channel would carry, contributes nothing measurable on a binary readout, on a graded readout with one-point resolution, or on exact probabilities. In Shoemaker’s terms, “I wrote that” is a report that requires identifying which thing is oneself, and it can misidentify: the model claims a filler turn it never generated and disclaims a word it produces on every fork the moment the word carries a user label. In Wegner’s terms, exclusivity operates on the Claude models as it does on people, and what it operates on is an inference, since the inference is defeated by an offered alternative author and untouched by the model’s own likelihood.

What the paragraph result adds. At paragraph length Claude Opus 5 and GPT-5.6-Sol carry a comparative signal that one-word answers cannot: they pick their own paragraph in forced choice, and on Claude Opus 5 the rival frame reads the same signal. A style signal at that length is what the third-person self-recognition literature predicts [12, 16]: given enough text, a model can recognise a register, its own included, as any reader who had seen its outputs could. Two facts keep the report on the reader’s side of the line: it does not appear under the neutral question, where the label decides alone (1,152 of 1,152 Yes), and on three of four comparisons it is indistinguishable from a length heuristic. But the one-word claim, that own probability contributes nothing, does not carry over to paragraphs, and we do not carry it. The hedged copy of a model’s own paragraph is both off-register and less probable, and every judge disowns it; the exact-probability arm shows that at 4B the rival-frame report rises with the paragraph’s log probability inside the model’s own register while no model at any scale can pick its own paragraph in forced choice. The one-word experiments hold register fixed and vary probability across 27 nats and find nothing; the paragraph experiments vary both together and find something. The honest statement is that the ownership report ignores own probability when probability is the only thing that varies, and that at paragraph length a signal appears which the production-model designs cannot separate from likelihood and the local design finds correlated with it.

Consequences. For agents that ask the model to check its own transcript, the answer is a transcript lookup: correct when the labels are correct, and




Figure 4: Paragraph length, clean harness (pilot 17c), twelve prompts, eight forks per cell. Left: rival-frame ownership of the judge’s own paragraph, of another model’s paragraph (pooled over the other two sources), and of the judge’s own paragraph with hedges prefixed to each sentence. Right: forced-choice accuracy with nothing planted, own paragraph against the other Claude model’s and against the other vendor’s, exact 95% intervals, dotted line at chance; Claude Haiku 4.5 counts are over the trials it answered (23 and 15 of 96).

no more. For model-based evaluation, self-preference driven by self-recognition [12] is a style effect that removing typography only partly attenuates. For introspection claims, first-person ownership on production models is not the place to look for a likelihood channel; the prefill-detection results of Lindsey [10] and Wang et al. [16] are consistent with ours if what is detected is style and preference mismatch rather than low own-probability. And for a self-model account of assistant behaviour, the self-model that emerges is an other-model pointed at the assistant turns, built from what any reader could see and wrong about the speaker’s idiosyncrasies (Claude Opus 5 says it would answer Prague and answers Lisbon) in exactly the way an outside reader with a generic prior would be.

6 Limitations

Own probabilities on the closed models are fork frequencies at 48 per prompt (log floor 1/96), not log-probs; the exact-probability arm covers only models up to 4B, where the design itself is fragile. The harness system prompt is part of the stimulus on both tools and is a plausible source of vendor differences. The two Claude judges share a developer; GPT-5.6-Sol was reached through one harness on one day. One-word answers are the cleanest case for likelihood and the weakest for style; the paragraph experiments trade the reverse way, carry a length confound that our design cannot remove, and confound register with likelihood, so the no-likelihood claim is a one-word claim. Every run up to the clean replications carried the experimenter’s instruction file in the context, which we disclose and which

the clean reruns address. All judgements were made by models, including the adversarial reviews and number checks; no human rated anything.

7 Reproducibility

Every row of every run is committed, every pre-registration is in the pilot document with its commit hash, and every table in this paper is regenerated by a named script from those rows. The prefill tools (`selfportrait/fork.py`, `selfportrait/codex_fork.py`) need only the two command-line tools and a subscription; the GPT-5.6-Sol arm has no dollar cost. The Claude runs were made on a Claude Max subscription through Claude Code; the per-call cost that `claude -p` reports is an API-equivalent figure, not a bill, and summed over every row in the repository (portrait experiments and contaminated runs included) it comes to about \$440. The local runs cost nothing beyond electricity.

Acknowledgements

Claude (Anthropic) was used throughout for code, adversarial review, independent recomputation of every reported figure, and drafting; the author is responsible for the design, the claims and the errors. The experiments were run on the author’s own accounts.

References

- [1] Thibaud Ardoin, Jonas Schäfer, and Gerhard Wunder. LLM self-recognition: Steering and retrieving activation signatures. *arXiv preprint arXiv:2606.06315*, 2026. URL <https://arxiv.org/abs/2606.06315>.
- [2] Xiaoyan Bai, Aryan Shrivastava, Ari Holtzman, and Chenhao Tan. Know thyself? on the incapability and implications of AI self-recognition. *arXiv preprint arXiv:2510.03399*, 2025. URL <https://arxiv.org/abs/2510.03399>.
- [3] Daryl J. Bem. Self-perception theory. In Leonard Berkowitz, editor, *Advances in Experimental Social Psychology*, volume 6, pages 1–62. Academic Press, New York, 1972.
- [4] Jan Betley, Xuchan Bao, Martín Soto, Anna Szttyber-Betley, James Chua, and Owain Evans. Tell me about yourself: LLMs are aware of their learned behaviors. *arXiv preprint arXiv:2501.11120*, 2025. URL <https://arxiv.org/abs/2501.11120>.
- [5] Felix J. Binder, James Chua, Tomek Korbak, Henry Sleight, John Hughes, Robert Long, Ethan Perez, Miles Turpin, and Owain Evans. Looking inward: Language models can learn about themselves by introspection. *arXiv preprint arXiv:2410.13787*, 2024. URL <https://arxiv.org/abs/2410.13787>.
- [6] Aaron Fanous, Jacob Goldberg, Ank A. Agarwal, Joanna Lin, Anson Zhou, Roxana Daneshjou, and Sanmi Koyejo. Syceval: Evaluating LLM sycophancy. *arXiv preprint arXiv:2502.08177*, 2025. URL <https://arxiv.org/abs/2502.08177>.
- [7] Patrick Haggard, Sam Clark, and Jeri Kalogeras. Voluntary action and conscious awareness. *Nature Neuroscience*, 5:382–385, 2002. doi: 10.1038/nn827.
- [8] Shomik Jain, Charlotte Park, Matt Viana, Ashia Wilson, and Dana Calacci. Interaction context often increases sycophancy in LLMs. *arXiv preprint arXiv:2509.12517*, 2025. URL <https://arxiv.org/abs/2509.12517>.
- [9] Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova Das-Sarma, Eli Tran-Johnson, Scott Johnston, Sheer El-Showk, Andy Jones, Nelson Elhage, Tristan Hume, Anna Chen, Yuntao Bai, Sam Bowman, Stanislav Fort, Deep Ganguli, Danny Hernandez, Josh Jacobson, Jackson Kernion, Shauna Kravec, Liane Lovitt, Kamal Ndousse, Catherine Olsson, Sam Ringer, Dario Amodei, Tom Brown, Jack Clark, Nicholas Joseph, Ben Mann, Sam McCandlish, Chris Olah, and Jared Kaplan. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022. URL <https://arxiv.org/abs/2207.05221>.
- [10] Jack Lindsey. Emergent introspective awareness in large language models. *arXiv preprint arXiv:2601.01828*, 2026. URL <https://arxiv.org/abs/2601.01828>.
- [11] Richard E. Nisbett and Timothy D. Wilson. Telling more than we can know: Verbal reports on mental processes. *Psychological Review*, 84(3):231–259, 1977. doi: 10.1037/0033-295X.84.3.231.
- [12] Arjun Panickssery, Samuel R. Bowman, and Shi Feng. LLM evaluators recognize and favor their own generations. *arXiv preprint arXiv:2404.13076*, 2024. URL <https://arxiv.org/abs/2404.13076>.
- [13] Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. Towards understanding sycophancy in language models. *arXiv preprint arXiv:2310.13548*, 2023. URL <https://arxiv.org/abs/2310.13548>.
- [14] Sydney Shoemaker. Self-reference and self-awareness. *Journal of Philosophy*, 65(19):555–567, 1968.
- [15] Jesse St. Amand, Callum Canavan, Sohaib Imran, Joseph Hewson, Aaron Lutz, Shi Feng, Puria Radmard, and Lennie Wells. Self-generated text recognition: Quality heuristics, cross-task transfer, and downstream bias in LLM evaluation. *arXiv preprint arXiv:2608.26159*, 2026. URL <https://arxiv.org/abs/2608.26159>.
- [16] Andy Wang, Parv Mahajan, David Demetri Africa, Alexandra Souly, Jordan Taylor, and Robert Kirk. Prefill awareness in large language models. *arXiv preprint arXiv:2606.12747*, 2026. URL <https://arxiv.org/abs/2606.12747>.
- [17] Daniel M. Wegner and Thalia Wheatley. Apparent mental causation: Sources of the experience of will. *American Psychologist*, 54(7):480–492, 1999. doi: 10.1037/0003-066X.54.7.480.